



Training runs to completion, and average episode reward does not improve over the entire budget


PPO — average reward per episode




PPO — value error (critic regression loss)




PPO — training-time explained variance (optimistically biased — see caption)



DQN — average reward per episode



DQN — value error (temporal-difference loss)



DQN — exploration rate (annealed)

